

Project 2

Statistical Patter Recognition

Girguis Sedky

1 Support Vector Machine

In this method, we use train a support vector machine (SVM) on the training FMNIST data, which consists of 60,000 images of 10 different clothing items and deploy the trained algorithm on the test data, which consists of 10,000 images. LIBSVM was used here. The dimension of the image vectors are reduced from $28 * 28 = 784$ to 50 dimensions using principal component analysis before the training procedure. We try three different Kernels, including a linear kernel, a third degree polynomial kernel, and a radial basis function kernel (RBF). Figure 1 shows a comparison of the accuracy of the SVM algorithm in classifying the test data with each kernel. All of the algorithms do relatively well. The SVM algorithm trained using the RBF kernel is the most accurate classifier.




Figure 1: Bayesian classifier

2 Convolutional Neural Network

In this section, we utilize a CNN to classify the test images. No PCA is performed on the data set that is used here. The network used is a variant on LeNet that is shown in Fig. 2. It consists of 7 layers and an alternating convolution and max pooling application from one layer to the next. ReLU is introduced in the final transition layer. MatCovNet toolbox was utilized to develop, train, and implement this neural network.

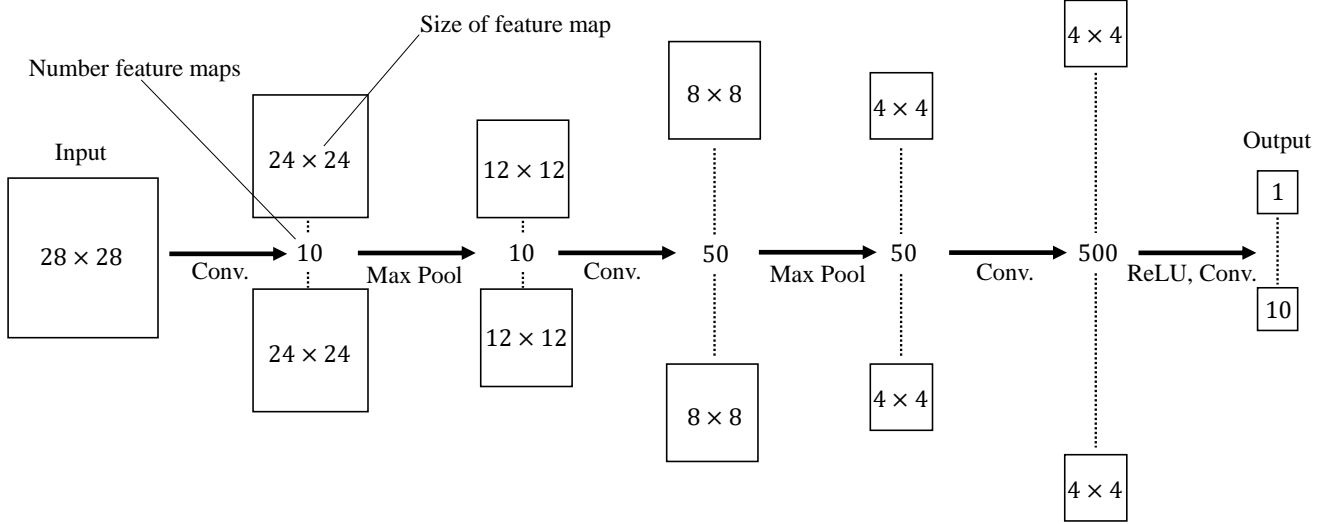


Figure 2: Convolutional Neural Network layout

The data is divided into batches of 100 images, and the network goes through the entire data set 15 times (15 epochs) during the gradient descent procedure to tune all the weights. Figure 3 shows the rate of descent of error as a function of the epoch number. The error rate has a steep drop during the first couple of passes through the data, and then it continues decreasing with a steadier pace. The trained CNN is able to classify the test images with 90.6% accuracy which is a higher accuracy compared to SVM. This improvement in accuracy can be attributed to the increased complexity of the classifier, large number of empirical weights that can be tuned in training, and non-linearity introduced by the ReLU operator.




Figure 3: The descent rate of error in classifying the training data

3 Summary

Table 1 shows the classification error of the different methods utilized in project 1 and 2. In Project 1, the most accurate method of classification was to use principal component analysis to reduce the dimensionality of the state vectors to 50 and then to apply Nearest-Neighbor (NN) classification to the transformed data.

In this project, both Support Vector Machines (SVM) and Convolutional Neural Networks (CNN) were superior in performance to all the methods used in project 1. CNN proves to be the best classifier. This can be attributed to the increased complexity of the classifier, large number of empirical weights that can be tuned in training, and non-linearity introduced by the ReLU operator.

Table 1: Summary of classification results

| Method | Percentage Error |
|------------------|------------------|
| SVM (RBF Kernel) | 12% |
| CNN | 9.39% |
| Bayes | 35% |
| NN | 15% |
| PCA, Bayes (50) | 20.5% |
| PCA, NN (150) | 14.7% |
| LDA, Bayes | 25.4% |
| LDA, NN | 28% |

Appendix

SVM

PCA

```
1 %% Reduce the dimension of the FMNIST data using PCA
2 clear; clc; close all;
3 % Girguis Sedky
4 %
5 %%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%
6 %% Load data
7 % Prototypes
8 images_train = loadMNISTImages('G:\My Drive\Classes&Books\
   StatisticalPatternRecognition\Project\Project2\Data\train-images-
   idx3-ubyte');
9 labels_train = loadMNISTLabels('G:\My Drive\Classes&Books\
   StatisticalPatternRecognition\Project\Project2\Data\train-labels-
   idx1-ubyte');
10 % Test data
11 images_test = loadMNISTImages('G:\My Drive\Classes&Books\
   StatisticalPatternRecognition\Project\Project2\Data\t10k-images-
   idx3-ubyte');
12 labels_test = loadMNISTLabels('G:\My Drive\Classes&Books\
   StatisticalPatternRecognition\Project\Project2\Data\t10k-labels-
   idx1-ubyte');
13
14 %% PCA
15 % Find the prinicpal components
16 coeff = pca(images_train');
17 % Number of dimensions kept
18 m = 50;
19
20 %% reduce the dimensionality of the system to m
21 images_train = coeff(:,1:m)'*images_train;
22 images_test = coeff(:,1:m)'*images_test;
23
24 %% Save the data
25 save('Data_PCA.mat','images_train','labels_train','images_test','
   labels_test');
```

Training script

```
1 %% Train SVM on reduced dimension FMNIST
2 clear; clc; close all;
3 % Girguis Sedky
4 %
   %%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%
5 %% Load the reduced data file
6 load('Data_PCA.mat');
7
8 %% Train SVM using different kernels
9 % linear kernel
10 model_linear = svmtrain(labels_train , images_train , '-t 0');
11 % polynomial kernel
12 model_poly = svmtrain(labels_train , images_train , '-t 1');
13 % radial basis function kernel
14 model_rbf = svmtrain(labels_train , images_train , '-t 2');
15
16 %% Save the model
17 save('SVM_Model.mat', 'model_linear', 'model_poly', 'model_rbf');
```

Testing script

```
1 %% Test SVM trained model on reduced dimension FMNIST test data
2 clear; clc; close all;
3 % Girguis Sedky
4 %
   %%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%
5
6 %% Load the reduced data file
7 load( 'Data_PCA.mat' );
8 load( 'SVM_Model.mat' );
9
10 %% Run SVM on test data
11 [predicted_label, accuracy_linear, dec_values] = svmpredict(
    labels_test, images_test', model_linear);
12 [predicted_label, accuracy_poly, dec_values] = svmpredict(labels_test,
    images_test', model_poly);
13 [predicted_label, accuracy_rbf, dec_values] = svmpredict(labels_test,
    images_test', model_rbf);
14
15 %% Save the accuracy of the SVM with different kernels
16 Performance = [accuracy_linear(1), accuracy_poly(1), accuracy_rbf(1)];
17 save( 'Performance.mat', 'Performance' )
```

CNN

Neural network initialization

```
1 % Create a LeNet
2 function net = initialize_FMNIST_CNN()
3
4 f=1/100 ;
5 net.layers = {} ;
6 net.layers{end+1} = struct('type', 'conv', ...
7                             'weights', {{f*randn(5,5,1,20), zeros(1,
8                                     20)}}}, ...
9                             'stride', 1, ...
10                             'pad', 0) ;
11 net.layers{end+1} = struct('type', 'pool', ...
12                             'method', 'max', ...
13                             'pool', [2 2], ...
14                             'stride', 2, ...
15                             'pad', 0) ;
16 net.layers{end+1} = struct('type', 'conv', ...
17                             'weights', {{f*randn(5,5,20,50), zeros(1,50)
18                                     }}, ...
19                             'stride', 1, ...
20                             'pad', 0) ;
21 net.layers{end+1} = struct('type', 'pool', ...
22                             'method', 'max', ...
23                             'pool', [2 2], ...
24                             'stride', 2, ...
25                             'pad', 0) ;
26 net.layers{end+1} = struct('type', 'conv', ...
27                             'weights', {{f*randn(4,4,50,500), zeros
28                                     (1,500)}}}, ...
29                             'stride', 1, ...
30                             'pad', 0) ;
31 net.layers{end+1} = struct('type', 'relu') ;
32 net.layers{end+1} = struct('type', 'conv', ...
33                             'weights', {{f*randn(1,1,500,10), zeros
34                                     (1,10)}}}, ...
35                             'stride', 1, ...
36                             'pad', 0) ;
37 net.layers{end+1} = struct('type', 'softmaxloss') ;
38
39 net = vl_simplenn_tidy(net) ;
```


Main training and testing script

```
1 % train and test FMNIST data via deep learning
2 % Girguis Sedky
3 function CNN_FMNIST_Script(varargin)
4 setup;
5 %%
6 %


---


7 % Part 4.1: prepare the data
8 %


---


9 %%
10 % Load training data set
11 % Prototypes
12 images_train = loadMNISTImages('G:\My Drive\Classes&Books\
    StatisticalPatternRecognition\Project\Project2\Data\train-images-
    idx3-ubyte');
13 labels_train = loadMNISTLabels('G:\My Drive\Classes&Books\
    StatisticalPatternRecognition\Project\Project2\Data\train-labels-
    idx1-ubyte');
14 labels_train=labels_train+1;
15 % reshape data so that you get a picture array
16 images_train = reshape(images_train, [28 28 60000]);
17 %%
18 %


---


19 % Part 4.2: initialize a CNN architecture
20 %


---


21 net = initialize_FMNIST_CNN();
22 vl_simplenn_display(net)
23
24 %%
25 %


---


26 % Part 4.3: train and evaluate the CNN
27 %


---


28 % Set the options for the stochastic gradient descent backpropagation
29 trainOpts.batchSize = 100; % The number of pictures that goes into
```

```

    one run of the descent
30 trainOpts.numEpochs = 15 ;      % Nummber of times the gradient
    procedure runs through all the data
31 trainOpts.continue = false ;    % If it were stopped, the descent
    continues where its left off
32 trainOpts.gpus = [] ;           % No GPU to be used
33 trainOpts.learningRate = 0.001 ; % Learning rate
34 trainOpts.expDir = 'NNData' ;   % directory where the options are
    saved
35 trainOpts = vl_argparse(trainOpts, varargin);
36
37 % Subtract out the mean of the image
38 im_mean = mean(images_train(:)) ;
39 images_train = images_train - im_mean ;
40
41 % Create the imdb structure that is apparently needed to run this
42 imdb.images.data = images_train;
43 imdb.images.label = labels_train';
44 imdb.images.id = [1:1:length(imdb.images.label)];
45 imdb.images.set = 1*ones(1,length(imdb.images.label));
46 % Call training function in MatConvNet
47 [net,info] = cnn_train(net, imdb, @getBatch, trainOpts);
48
49
50 % Save the result for later use
51 net.layers(end) = [] ;
52 net.imageMean = im_mean ;
53 save('NNData/CNN.mat', '-struct', 'net') ;
54
55 %%
56 %

```

```

57 % Part 4.5: apply the model
58 %

```

```

59 clear;
60 % Load the CNN learned before
61 net = load('NNData/CNN.mat') ;
62
63 % Load test data set
64 % Prototypes
65 images_test = loadMNISTImages('G:\My Drive\Classes&Books\
    StatisticalPatternRecognition\Project\Project2\Data\t10k-images-
    idx3-ubyte');
66 labels_test = loadMNISTLabels('G:\My Drive\Classes&Books\

```

```

        StatisticalPatternRecognition\Project\Project2\Data\t10k-labels-
        idx1-ubyte');
67 labels_test=labels_test+1;
68 % reshape data so that you get a picture array
69 images_test = reshape(images_test, [28 28 10000]);
70
71 % subtract mean and ,multiply by that weird factor
72 images_test = 256 * (images_test- net.imageMean) ;
73
74 % Image to be tested
75 for ii=1:length(images_test)
76
77 % Apply the CNN to the test image
78 res = vl_simplenn(net, images_test(:,:,ii)) ;
79
80 %
81 scores = squeeze(gather(res(end).x)) ;
82 [~, best] = max(scores);
83 labels_est(ii) = best;
84 end
85
86 labels_est = labels_est';
87 % Test accuracy of classifying algorithm
88 a = find(labels_est==labels_test);
89 Performance = length(a)*100/length(labels_test)
90 save('Performance.mat','Performance');
91 % -----
92
93 function [im, labels] = getBatch(imdb, batch)
94 %% Custo, function to create the batches that go into the descent
    procedure
95 % -----
96 im = imdb.images.data(:,:,batch) ;
97 im = 256 * reshape(im, 28, 28, 1, []) ;
98 labels = imdb.images.label(1,batch) ;

```